

Lecture 07

BFS / DFS / Shortest paths problem

Spring 2023

Zhihua Jiang

Graphs I: Breadth First Search

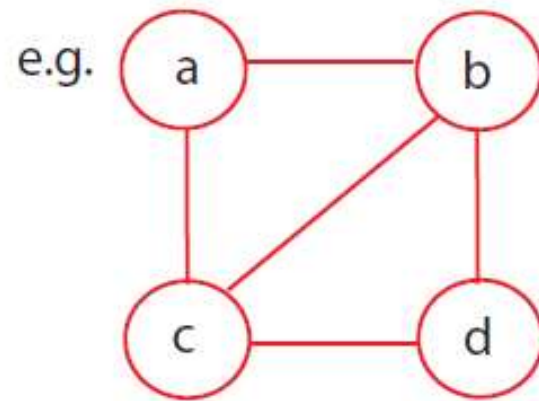
Lecture Overview

- Applications of Graph Search
- Graph Representations
- Breadth-First Search

Recall:

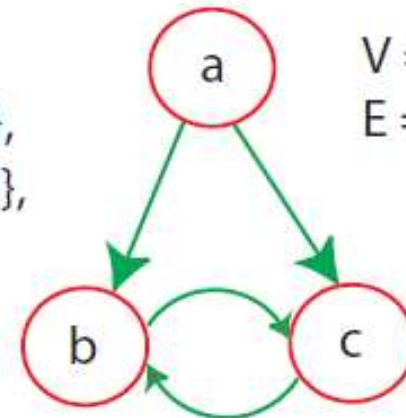
Graph $G = (V, E)$

- V = set of vertices (**arbitrary labels**)
- E = set of edges i.e. vertex pairs (v, w)
 - ordered pair \implies directed edge of graph
 - unordered pair \implies undirected



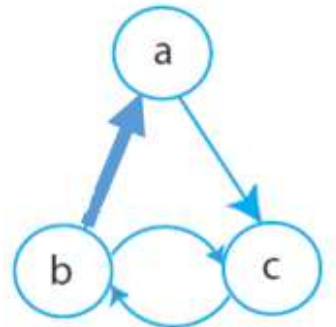
UNDIRECTED

$V = \{a, b, c, d\}$
 $E = \{\{a, b\}, \{a, c\},$
 $\{b, c\}, \{b, d\},$
 $\{c, d\}\}$



DIRECTED

$V = \{a, b, c\}$
 $E = \{(a, c), (b, c),$
 $(c, b), (b, a)\}$



Graph Representations: (data structures)

1

Adjacency lists:

Array Adj of $|V|$ linked lists

- for each vertex $u \in V$, $Adj[u]$ stores u 's neighbors, i.e., $\{v \in V \mid (u, v) \in E\}$. (u, v) are just outgoing edges if directed. (See [Fig. 2](#) for an example.)

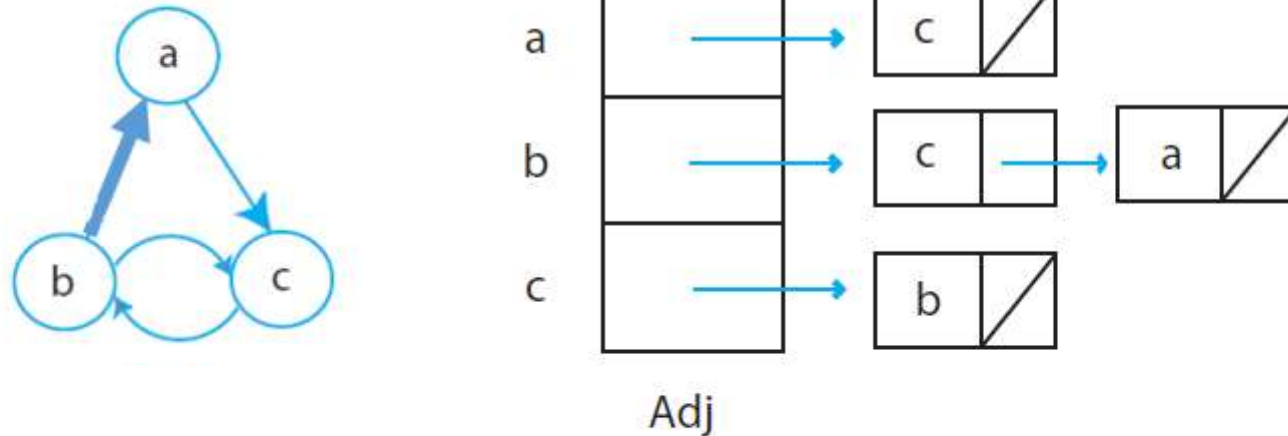
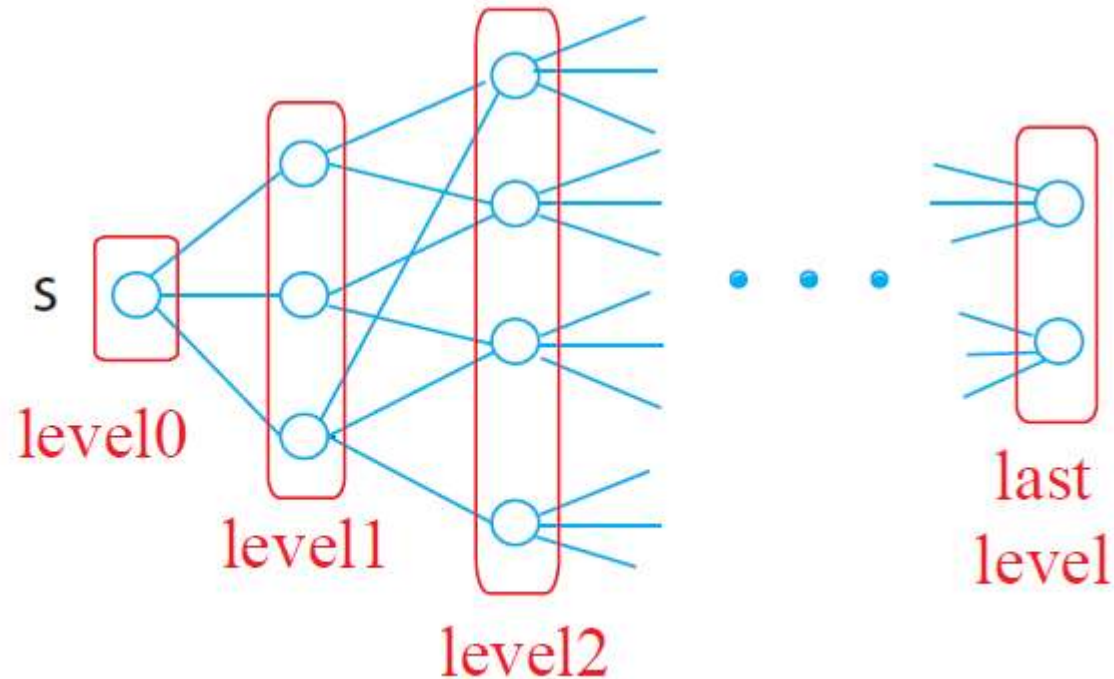


Figure 2: Adjacency List Representation: Space $\Theta(V + E)$

Breadth-First Search

Explore graph level by level from s

- level 0 = $\{s\}$
- level i = vertices reachable by path of i edges but not fewer



- build level $i > 0$ from level $i - 1$ by trying all outgoing edges, but ignoring vertices from previous levels

Breadth-First-Search Algorithm

```
BFS (V,Adj,s):  
    level = { s: 0 }  
    parent = { s : None }  
    i = 1  
    frontier = [s]                                # previous level, i - 1  
    while frontier:  
        next = [ ]                                # next level, i  
        for u in frontier:  
            for v in Adj[u]:  
                if v not in level:                 # not yet seen  
                    level[v] = i                   # = level[u] + 1  
                    parent[v] = u  
                    next.append(v)  
        frontier = next  
    i += 1
```

Analysis:

- vertex V enters next (& then frontier)
only once (because $\text{level}[v]$ then set)

base case: $v = s$

- $\implies \text{Adj}[v]$ looped through only once

$$\text{time} = \sum_{v \in V} |\text{Adj}[v]| = \begin{cases} |E| & \text{for directed graphs} \\ 2|E| & \text{for undirected graphs} \end{cases}$$

- $\implies O(E)$ time
- $O(V + E)$ (“**LINEAR TIME**”) to also list vertices unreachable from v (those still not assigned level)

Shortest Paths:

- for every vertex v , fewest edges to get from s to v is

$$\begin{cases} \text{level}[v] & \text{if } v \text{ assigned level} \\ \infty & \text{else (no path)} \end{cases}$$

- parent pointers form shortest-path tree = union of such a shortest path for each v
 \implies to find shortest path, take v , $\text{parent}[v]$, $\text{parent}[\text{parent}[v]]$, etc., until s (or None)

Graphs II: Depth-First Search

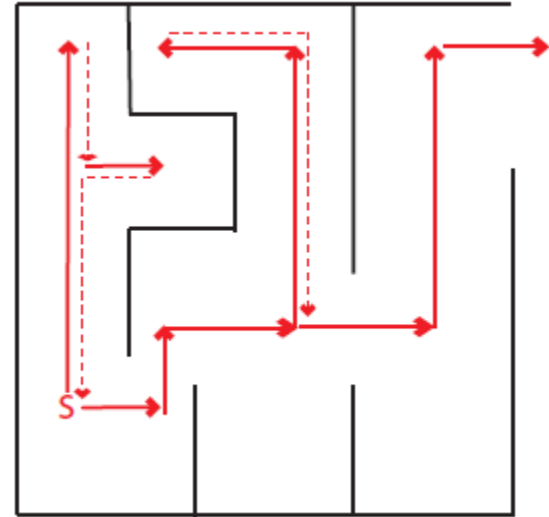
Lecture Overview

- Depth-First Search
- Edge Classification
- Cycle Testing
- Topological Sort

Depth-First Search (DFS)

This is like exploring a maze.

- follow path until you get stuck
- backtrack along breadcrumbs until reach unexplored neighbor
- recursively explore
- careful not to repeat a vertex



```

parent = {s: None}

DFS-visit (V, Adj, s):
  start → for v in Adj [s]:
    v      if v not in parent:
            parent [v] = s
            DFS-visit (V, Adj, v)
  finish → v

DFS (V, Adj)
  parent = { }
  for s in V:
    if s not in parent:
      parent [s] = None
      DFS-visit (V, Adj, s)

```

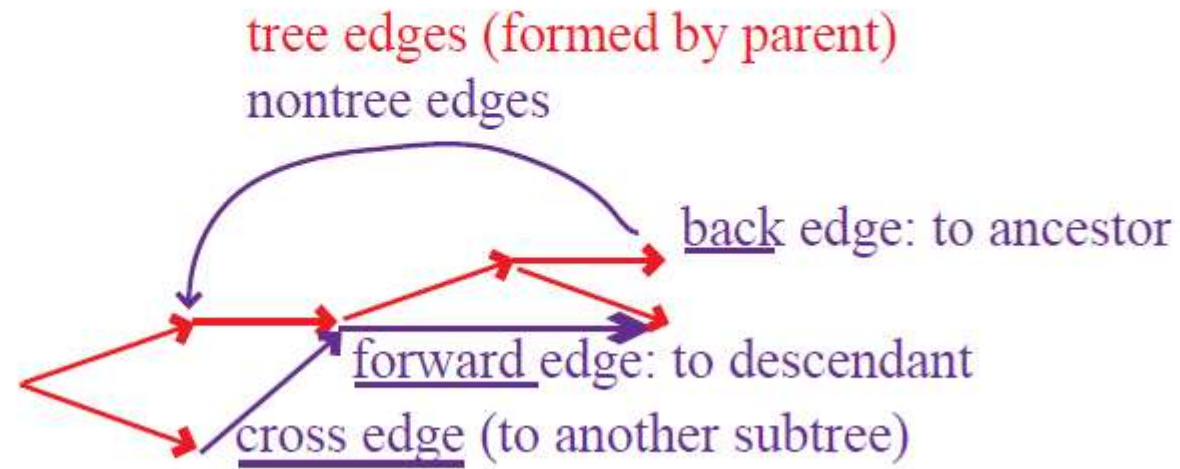
search from start vertex s
(only see stuff reachable from s)

explore entire graph

Analysis

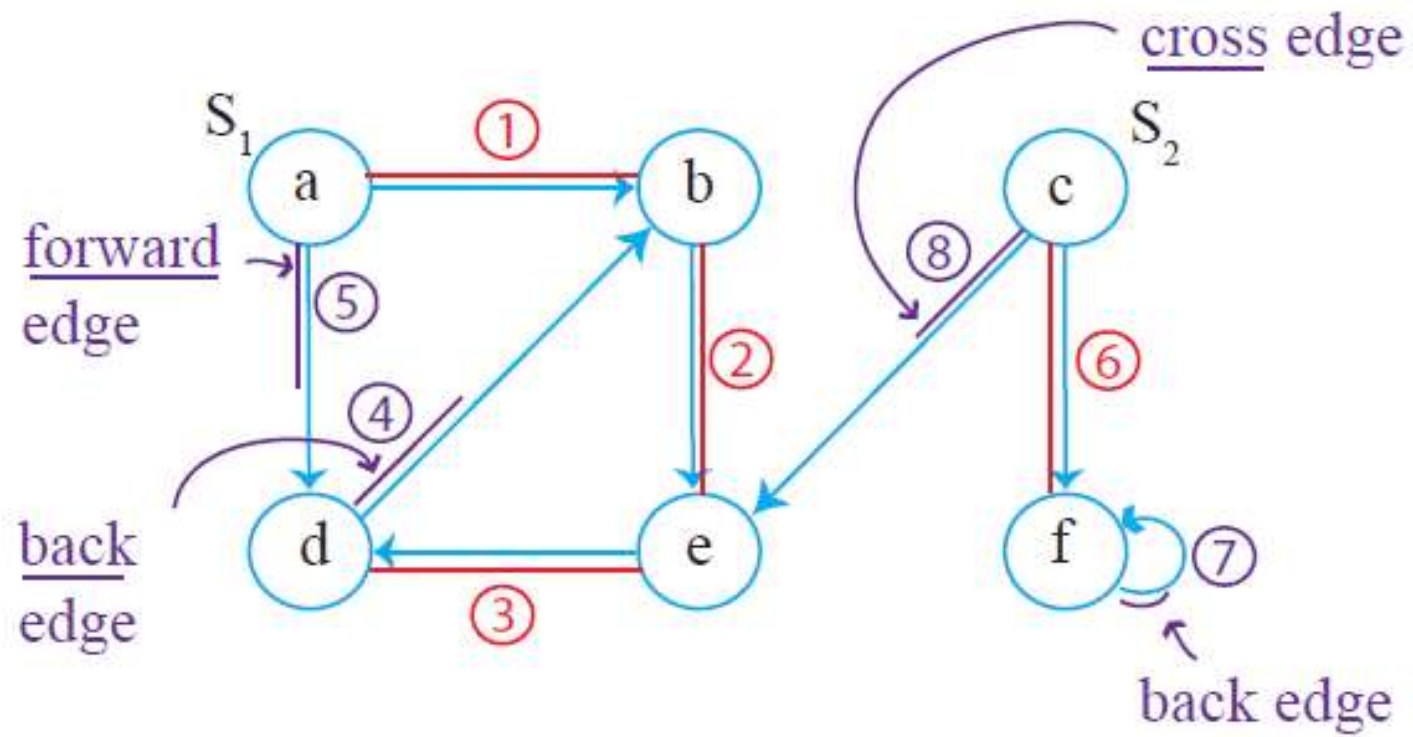
- DFS-visit gets called with a vertex s only once (because then $\text{parent}[s]$ set)
 \implies time in DFS-visit = $\sum_{s \in V} |\text{Adj}[s]| = O(E)$
- DFS outer loop adds just $O(V)$
 $\implies O(V + E)$ time (linear time)

Edge Classification



- to compute this classification (**back or not**), mark nodes for duration they are “on the stack”
- only tree and back edges in undirected graph

Example



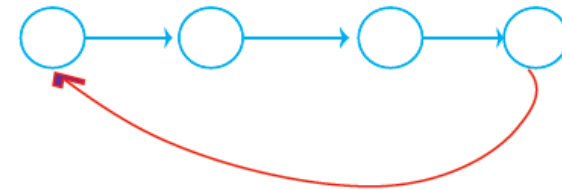
Cycle Detection

Graph G has a cycle \Leftrightarrow DFS has a back edge

- before visit to v_i finishes,
will visit v_{i+1} (& finish):
will consider edge (v_i, v_{i+1})
 \Rightarrow visit v_{i+1} now or already did
- \Rightarrow before visit to v_0 finishes,
will visit v_k (& didn't before)
- \Rightarrow before visit to v_k (or v_0) finishes,
will see (v_k, v_0) as back edge

(\Leftarrow)

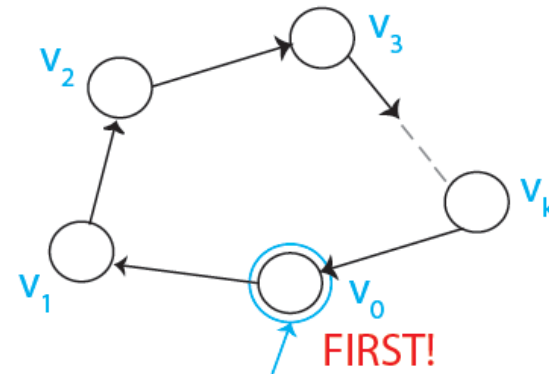
tree edges



is a cycle

back edge: to tree ancestor

(\Rightarrow) consider first visit to cycle:



Job scheduling

Given Directed Acyclic Graph (DAG), where vertices represent tasks & edges represent dependencies, order tasks without violating dependencies

Source = vertex with no incoming edges
= schedulable at beginning

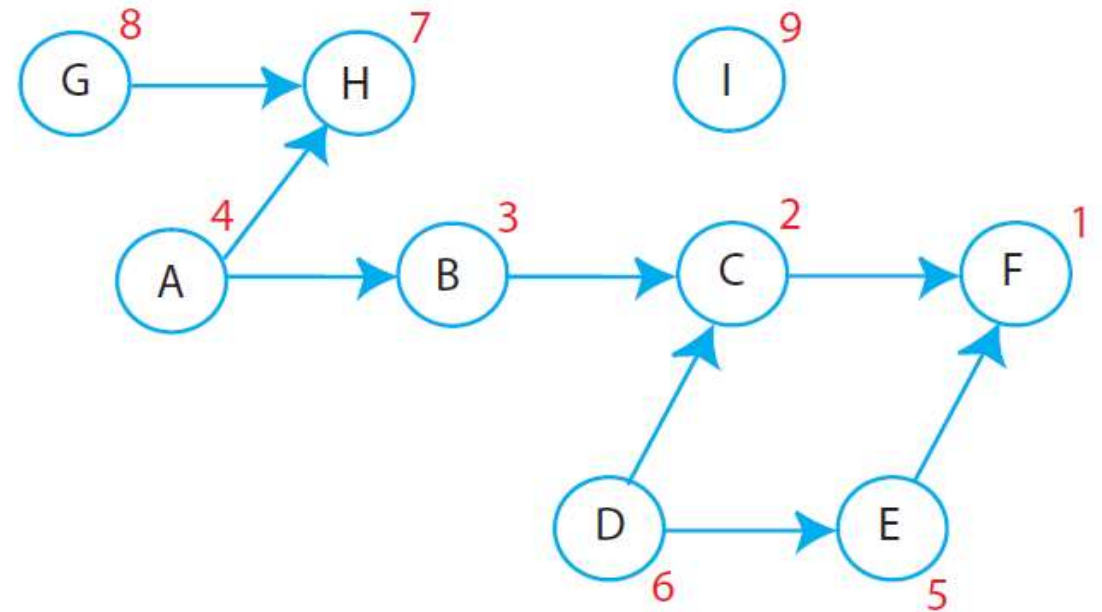


Figure 6: Dependence Graph: DFS Finishing Times

Topological Sort

Reverse of DFS finishing times (time at which DFS-Visit(v) finishes)

Correctness

For any edge (u, v) — u ordered before v , i.e., v finished before u



DFS-Visit(v)

...

order.append(v)

order.reverse()

Shortest Paths I: Intro

Lecture Overview

- Weighted Graphs
- General Approach
- Negative Edges
- Optimal Substructure

Motivation:

Shortest way to drive from A to B Google maps “get directions”

Formulation: Problem on a weighted graph $G(V, E)$ $W : E \rightarrow \mathfrak{R}$

Two algorithms: Dijkstra $O(V \lg V + E)$ assumes non-negative edge weights

Bellman Ford $O(VE)$ is a general algorithm

path $p = \langle v_0, v_1, \dots, v_k \rangle$

$(v_i, v_{i+1}) \in E$ for $0 \leq i < k$

$$w(p) = \sum_{i=0}^{k-1} w(v_i, v_{i+1})$$

Weighted Graphs:

Notation:

$v_0 \xrightarrow{p} v_k$ means p is a path from v_0 to v_k . (v_0) is a path from v_0 to v_0 of weight 0.

Definition:

Shortest path weight from u to v as

$$\delta(u, v) = \begin{cases} \min \left\{ w(p) : u \xrightarrow{p} v \right\} & \text{if } \exists \text{ any such path} \\ \infty & \text{otherwise } (v \text{ unreachable from } u) \end{cases}$$

Single Source Shortest Paths:

Given $G = (V, E)$, w and a source vertex S , find $\delta(S, V)$ [and the best path] from S to each $v \in V$.

Data structures:

$$\begin{aligned}d[v] &= \text{value inside circle} \\&= \left\{ \begin{array}{ll} 0 & \text{if } v = s \\ \infty & \text{otherwise} \end{array} \right\} \Leftarrow \text{initially} \\&= \delta(s, v) \Leftarrow \text{at end} \\d[v] &\geq \delta(s, v) \quad \text{at all times}\end{aligned}$$

$d[v]$ decreases as we find better paths to v

$\Pi[v]$ = predecessor on best path to v , $\Pi[s] = \text{NIL}$

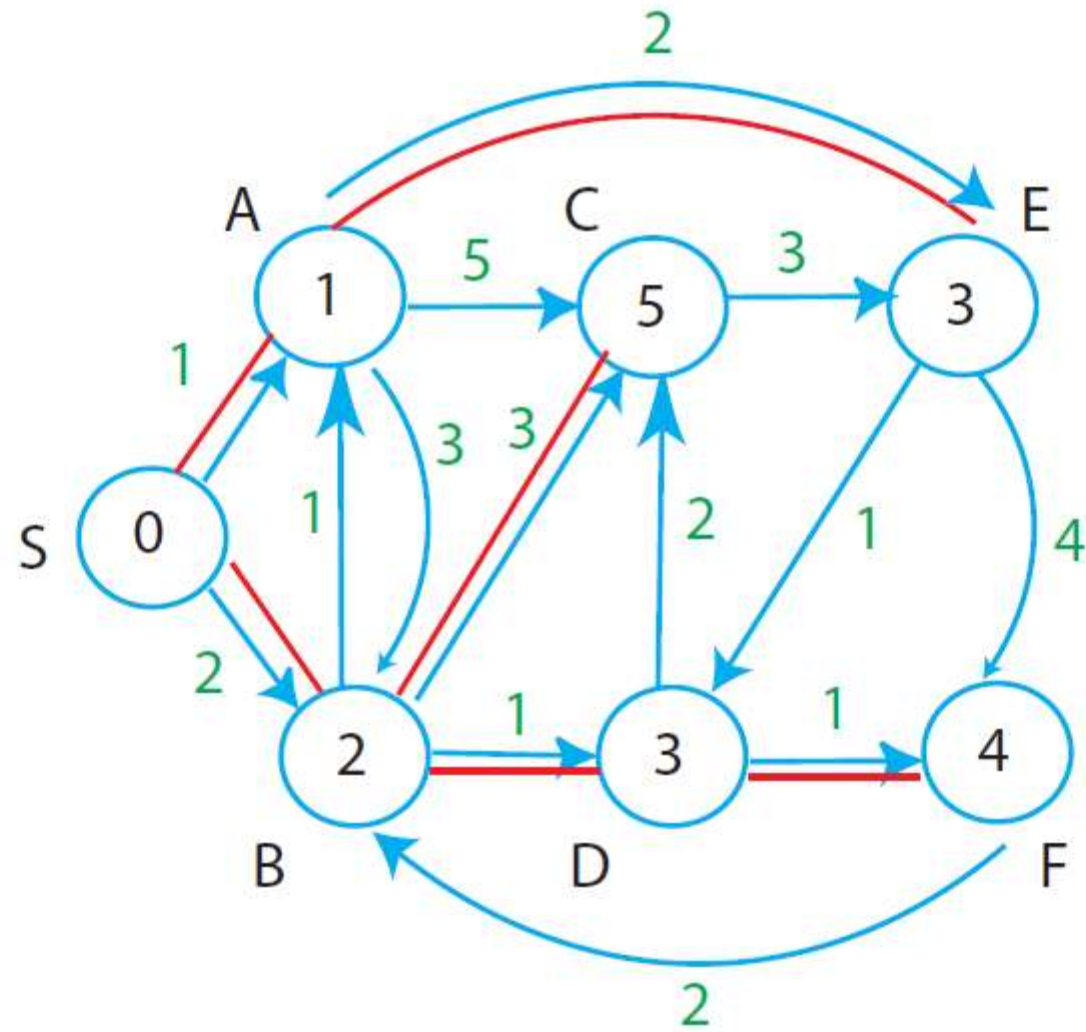
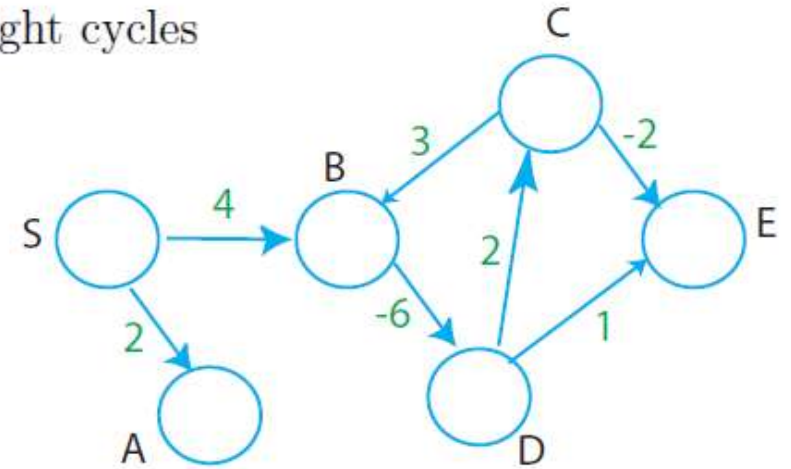


Figure 1: Shortest Path Example: Bold edges give predecessor Π relationships

Negative-Weight Edges:

- Natural in some applications (e.g., logarithms used for weights)
- Some algorithms disallow negative weight edges (e.g., Dijkstra)
- If you have negative weight edges, you might also have negative weight cycles
 \implies may make certain shortest paths undefined!



$B \rightarrow D \rightarrow C \rightarrow B$ (origin) has weight $-6 + 2 + 3 = -1 < 0$!

Shortest path $S \rightarrow C$ (or B, D, E) is undefined. Can go around $B \rightarrow D \rightarrow C$ as many times as you like

Shortest path $S \rightarrow A$ is defined and has weight 2

If negative weight edges are present, s.p. algorithm should find negative weight cycles (e.g., Bellman Ford)

General structure of S.P. Algorithms (no negative cycles)

```

Initialize:      for  $v \in V$ :  $d[v] \leftarrow \infty$ 
                   $\pi[v] \leftarrow \text{NIL}$ 
                   $d[S] \leftarrow 0$ 

Main:           repeat
                  select edge  $(u, v)$  [somehow]
                  [ if  $d[v] > d[u] + w(u, v)$  :
                     $d[v] \leftarrow d[u] + w(u, v)$ 
                     $\pi[v] \leftarrow u$ 
                  ]
                  until all edges have  $d[v] \leq d[u] + w(u, v)$ 

```

Optimal Substructure:

Theorem: Subpaths of shortest paths are shortest paths

Let $p = \langle v_0, v_1, \dots, v_k \rangle$ be a shortest path

Let $p_{ij} = \langle v_i, v_{i+1}, \dots, v_j \rangle \quad 0 \leq i \leq j \leq k$

Then p_{ij} is a shortest path.

Proof: $p =$

$$\begin{array}{ccccccc} & p_{0,i} & & p_{ij} & & p_{jk} & \\ v_0 & \rightarrow & v_i & \rightarrow & v_j & \rightarrow & v_k \\ & & & \rightarrow & & & \\ & & & p'_{ij} & & & \end{array}$$

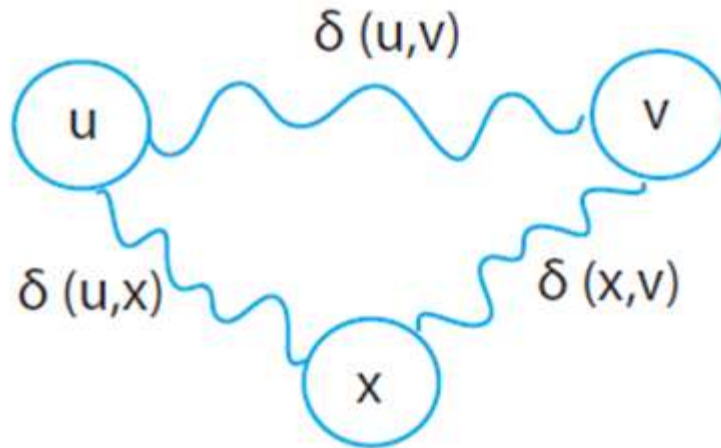
If p'_{ij} is shorter than p_{ij} , cut out p_{ij} and replace with p'_{ij} ; result is shorter than p .

Contradiction.

Triangle Inequality:

Theorem: For all $u, v, x \in X$, we have

$$\delta(u, v) \leq \delta(u, x) + \delta(x, v)$$



Shortest Paths II - Dijkstra

Lecture Overview

- Review
- Shortest paths in DAGs
- Shortest paths in graphs without negative edges
- Dijkstra's Algorithm

Relaxation is Safe

Lemma: The relaxation algorithm maintains the invariant that $d[v] \geq \delta(s, v)$ for all $v \in V$.

Proof: By induction on the number of steps.

Consider $RELAX(u, v, w)$. By induction $d[u] \geq \delta(s, u)$. By the triangle inequality, $\delta(s, v) \leq \delta(s, u) + \delta(u, v)$. This means that $\delta(s, v) \leq d[u] + w(u, v)$, since $d[u] \geq \delta(s, u)$ and $w(u, v) \geq \delta(u, v)$. So setting $d[v] = d[u] + w(u, v)$ is safe. \square

DAGs:

Can't have negative cycles because there are no cycles!

1. Topologically sort the DAG. Path from u to v implies that u is before v in the linear ordering.
2. One pass over vertices in topologically sorted order relaxing each edge that leaves each vertex.

$\Theta(V + E)$ time

Can deal with negative edges

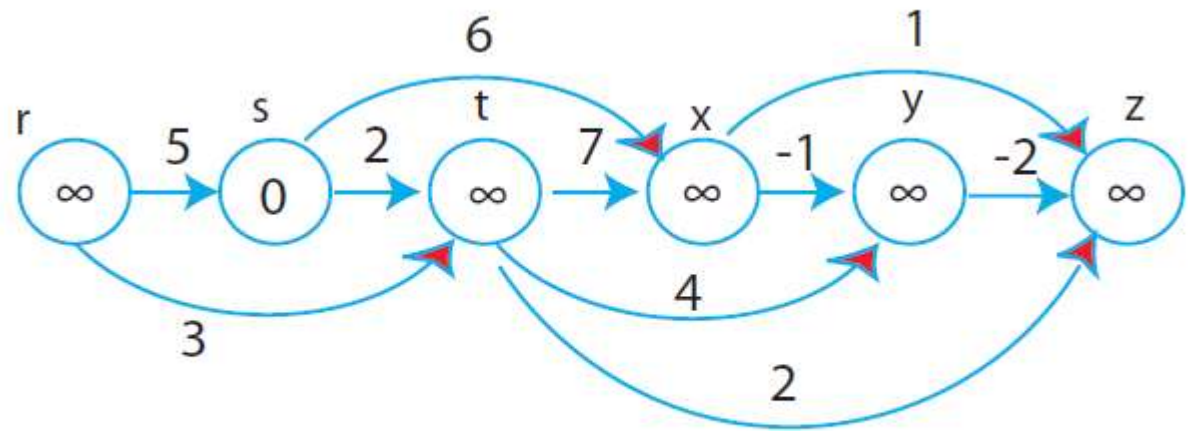
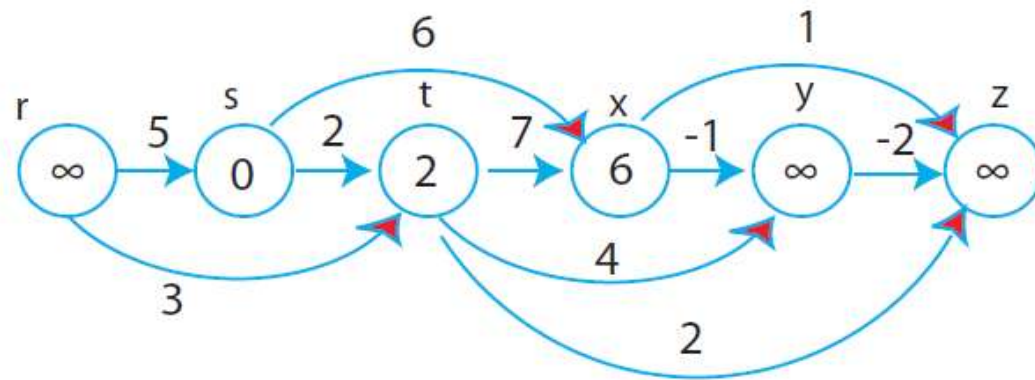


Figure 1: Shortest Path using Topological Sort.

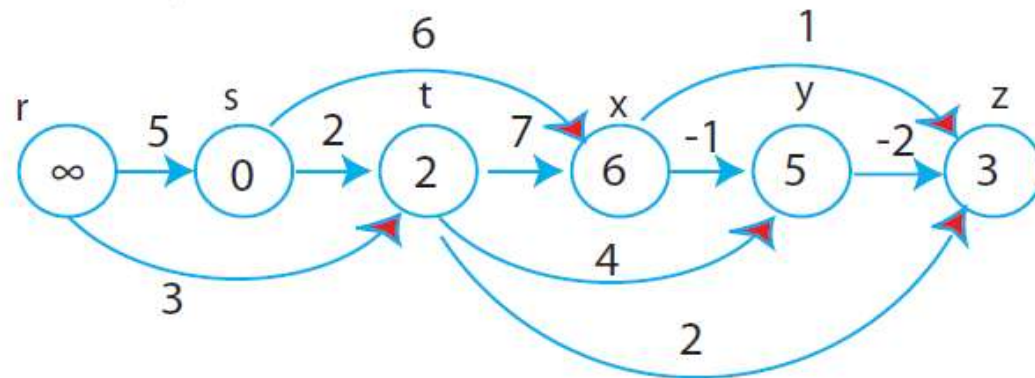
Vertices sorted left to right in topological order

Process r : stays ∞ . All vertices to the left of s will be ∞ by definition

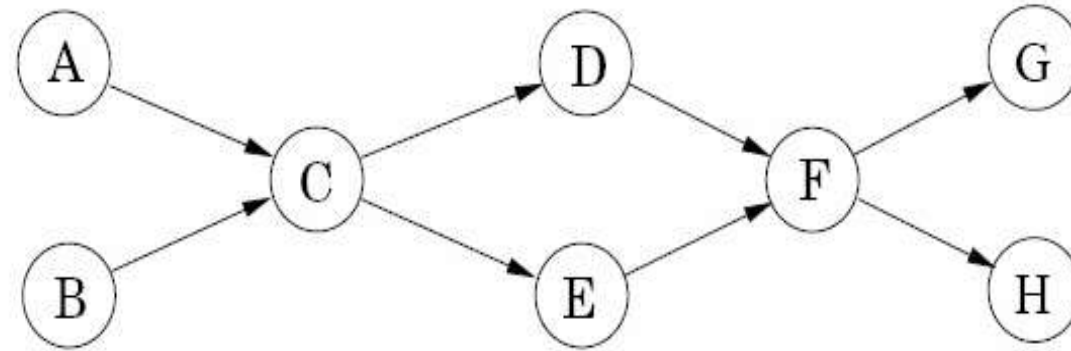
Process s : $t : \infty \rightarrow 2$ $x : \infty \rightarrow 6$



process t, x, y

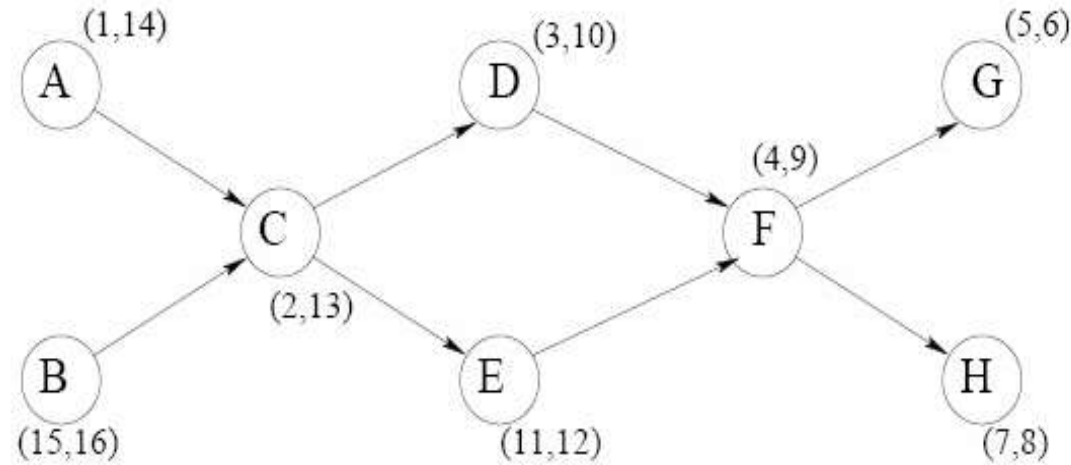


Run the DFS-based topological ordering algorithm on the following graph. Whenever you have a choice of vertices to explore, always pick the one that is alphabetically first.



- (a) Indicate the pre and post numbers of the nodes.
- (b) What are the sources and sinks of the graph?
- (c) What topological ordering is found by the algorithm?
- (d) How many topological orderings does this graph have?

(a) The figure below shows the **pre** and **post** times in parentheses.



(b) The vertices A, B are sources and G, H are sinks.

(c) Since the algorithm outputs vertices in decreasing order of **post** numbers, the ordering given is B, A, C, E, D, F, H, G .

(d) Any ordering of the graph must be of the form $\{A, B\}, C, \{D, E\}, F, \{G, H\}$, where $\{A, B\}$ indicates A and B may be in any order within these two places. Hence the total number of orderings is $2^3 = 8$.

procedure explore(G, v)

Input: $G = (V, E)$ is a graph; $v \in V$

Output: $\text{visited}(u)$ is set to true for all nodes u reachable from v

$\text{visited}(v) = \text{true}$

$\text{previsit}(v)$

for each edge $(v, u) \in E$:

 if not $\text{visited}(u)$: $\text{explore}(u)$

$\text{postvisit}(v)$

procedure dfs(G)

for all $v \in V$:

$\text{visited}(v) = \text{false}$

for all $v \in V$:

 if not $\text{visited}(v)$: $\text{explore}(v)$

procedure previsit(v)

$\text{pre}[v] = \text{clock}$

$\text{clock} = \text{clock} + 1$

procedure postvisit(v)

$\text{post}[v] = \text{clock}$

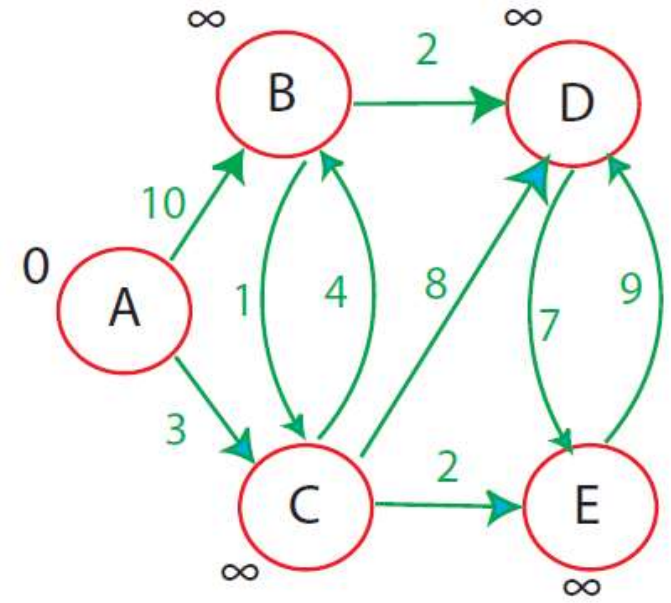
$\text{clock} = \text{clock} + 1$

Dijkstra's Algorithm

For each edge $(u, v) \in E$, assume $w(u, v) \geq 0$, maintain a set S of vertices whose final shortest path weights have been determined. Repeatedly select $u \in V - S$ with minimum shortest path estimate, add u to S , relax all edges out of u .

Pseudo-code

```
Dijkstra ( $G, W, s$ )    //uses priority queue  $Q$ 
    Initialize ( $G, s$ )
     $S \leftarrow \phi$ 
     $Q \leftarrow V[G]$     //Insert into  $Q$ 
    while  $Q \neq \phi$ 
        do  $u \leftarrow \text{EXTRACT-MIN}(Q)$     //deletes  $u$  from  $Q$ 
         $S = S \cup \{u\}$ 
        for each vertex  $v \in \text{Adj}[u]$ 
            do RELAX ( $u, v, w$ )    ← this is an implicit DECREASE_KEY operation
```



$S = \{ \}$	{ A B C D E }	=	Q	
$S = \{ A \}$	0 ∞ ∞ ∞ ∞			
$S = \{ A, C \}$	0 10 3 ∞ ∞	←		after relaxing edges from A
$S = \{ A, C \}$	0 7 3 11 5	←		after relaxing edges from C
$S = \{ A, C, E \}$	0 7 3 11 5			
$S = \{ A, C, E, B \}$	0 7 3 9 5	←		after relaxing edges from B

Strategy: Dijkstra is a greedy algorithm: choose closest vertex in $V - S$ to add to set S .

Correctness: We know relaxation is safe. The key observation is that each time a vertex u is added to set S , we have $d[u] = \delta(s, u)$.

Dijkstra Complexity

$\Theta(v)$ inserts into priority queue

$\Theta(v)$ EXTRACT_MIN operations

$\Theta(E)$ DECREASE_KEY operations

Priority queue implementations

- The running time of Dijkstra's algorithm depends heavily on the priority queue implementation.

Implementation	deletemin	insert/ decreasekey	$ V \times \text{deletemin} + (V + E) \times \text{insert}$
Array	$O(V)$	$O(1)$	$O(V ^2)$
Binary heap	$O(\log V)$	$O(\log V)$	$O((V + E) \log V)$
d -ary heap	$O(\frac{d \log V }{\log d})$	$O(\frac{\log V }{\log d})$	$O((V \cdot d + E) \frac{\log V }{\log d})$
Fibonacci heap	$O(\log V)$	$O(1)$ (amortized)	$O(V \log V + E)$

d -ary heap

- Identical to a binary heap, except that nodes have d children.
- The height of the tree: $\Theta(\log_d |V|) = \Theta(\log |V| / \log d)$
- Insert/DecreaseKey: speeded up, $\Theta(\log_d |V|)$
- DeleteMin: take a longer time, $\Theta(d \log_d |V|)$
- Total: $|V|\text{DeleteMin} + |V|\text{Insert} + |E|\text{DecreaseKey}$
$$= \Theta(|V| \cdot d \log_d |V| + (|V| + |E|) \cdot \log_d |V|)$$
$$= \Theta((|V| d + |E|) \log_d |V|)$$

Shortest Paths III: Bellman-Ford

Lecture Overview

- Review: Notation
- Generic S.P. Algorithm
- Bellman-Ford Algorithm
 - Analysis
 - Correctness

Bellman-Ford(G, W, s)

```
Initialize ()  
for  $i = 1$  to  $|V| - 1$   
    for each edge  $(u, v) \in E$ :  
        Relax( $u, v$ )  
for each edge  $(u, v) \in E$   
    do if  $d[v] > d[u] + w(u, v)$   
        then report a negative-weight cycle exists
```

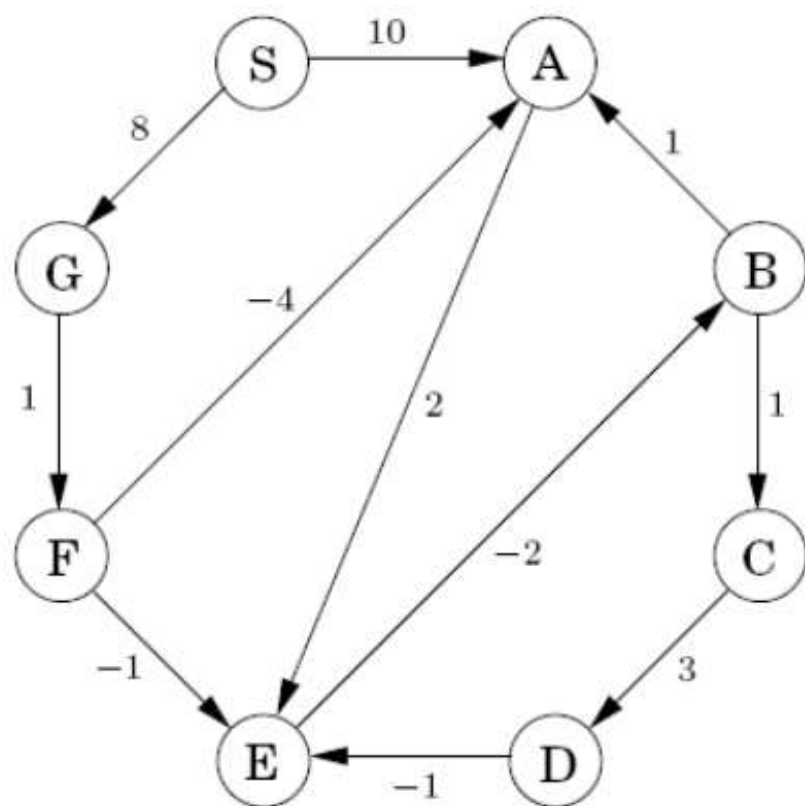
At the end, $d[v] = \delta(s, v)$, if no negative-weight cycles.

Corollary

If a value $d[v]$ fails to converge after $|V| - 1$ passes, there exists a negative-weight cycle reachable from s .

Theorem:

If $G = (V, E)$ contains no negative weight cycles, then after Bellman-Ford executes $d[v] = \delta(s, v)$ for all $v \in V$.



	Iteration							
Node	0	1	2	3	4	5	6	7
S	0	0	0	0	0	0	0	0
A	∞	10	10	5	5	5	5	5
B	∞	∞	∞	10	6	5	5	5
C	∞	∞	∞	∞	11	7	6	6
D	∞	∞	∞	∞	∞	14	10	9
E	∞	∞	12	8	7	7	7	7
F	∞	∞	9	9	9	9	9	9
G	∞	8	8	8	8	8	8	8

Shortest Paths IV - Speeding up Dijkstra

Lecture Overview

- Single-source single-target Dijkstra
- Bidirectional search
- Goal directed search - potentials and landmarks

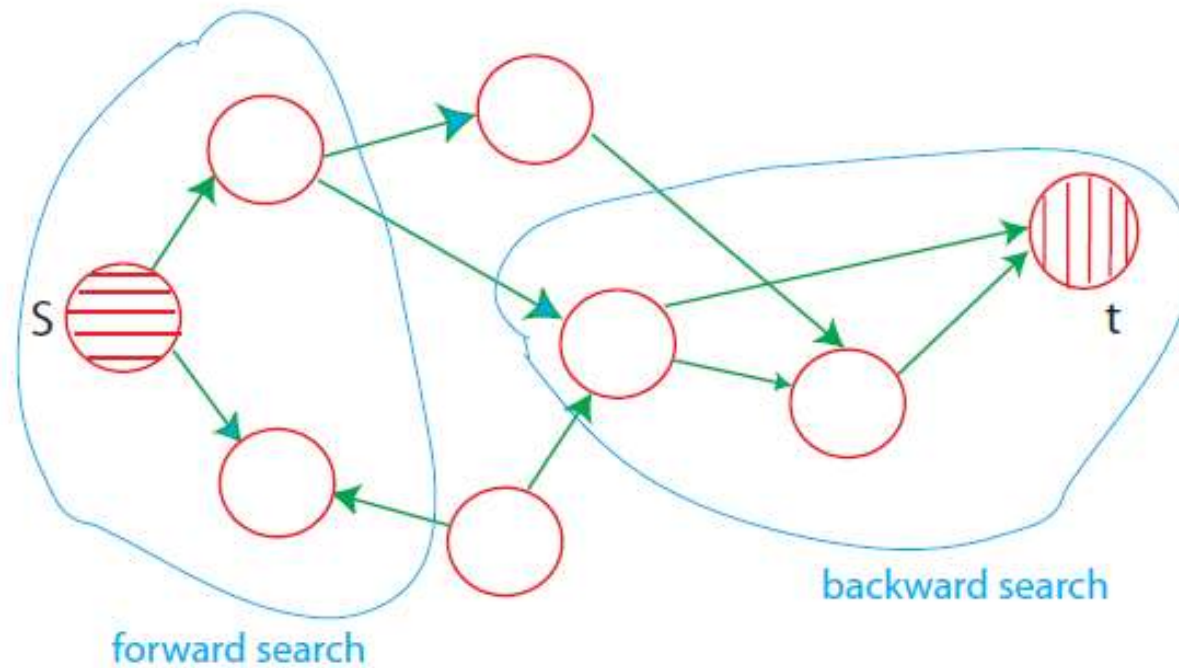
DIJKSTRA single-source, single-target

```
Initialize()
 $Q \leftarrow V[G]$ 
while  $Q \neq \phi$ 
    do  $u \leftarrow \text{EXTRACT\_MIN}(Q)$  (stop if  $u = t$ !)
    for each vertex  $v \in \text{Adj}[u]$ 
        do  $\text{RELAX}(u, v, w)$ 
```

Observation: If only shortest path from s to t is required, stop when t is removed from Q , i.e., when $u = t$

Bi-Directional Search

Note: Speedup techniques covered here do not change worst-case behavior, but reduce the number of visited vertices in practice.



Bi-D Search

Alternate forward search from s
backward search from t
(follow edges backward)
 $d_f(u)$ distances for forward search
 $d_b(u)$ distances for backward search

Algorithm terminates when some vertex w has been processed, i.e., deleted from the queue of both searches, Q_f and Q_b

Subtlety: After search terminates, find node x with minimum value of $d_f(x) + d_b(x)$.
 x may not be the vertex w that caused termination as in example to the left!

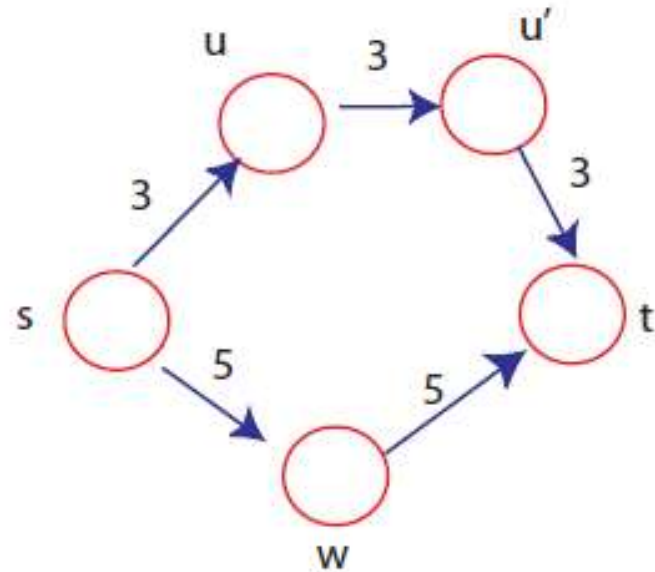
Find shortest path from s to x using Π_f and shortest path backwards from t to x using Π_b . *Note:* x will have been deleted from either Q_f or Q_b or both.

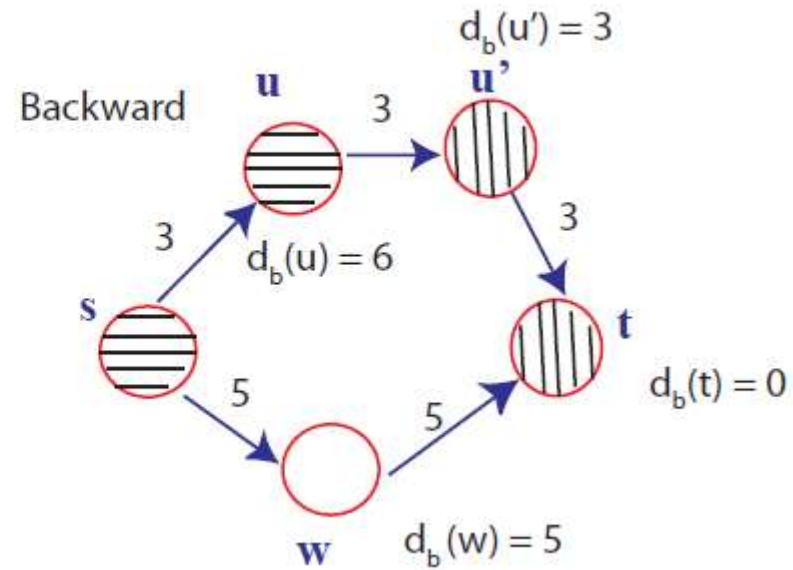
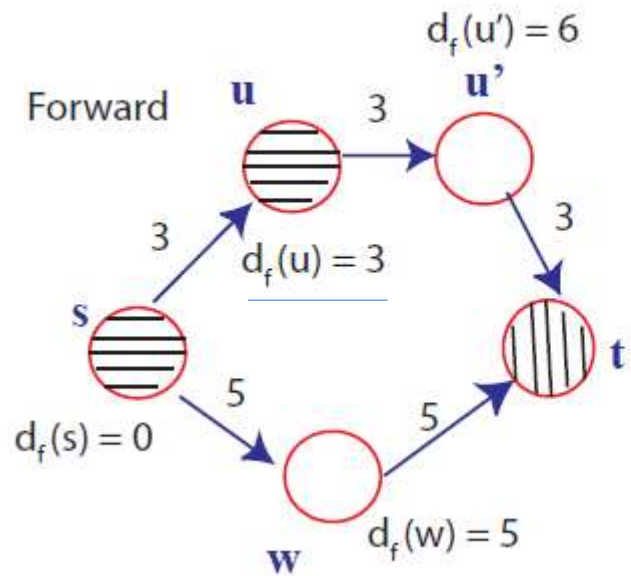
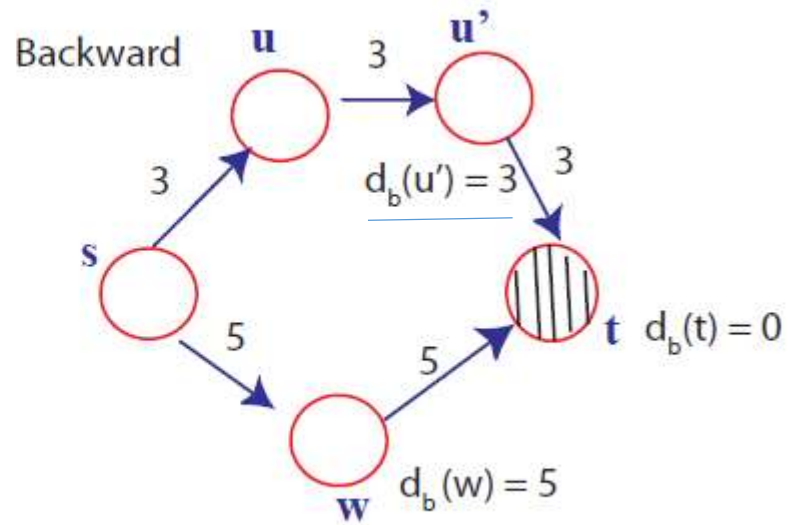
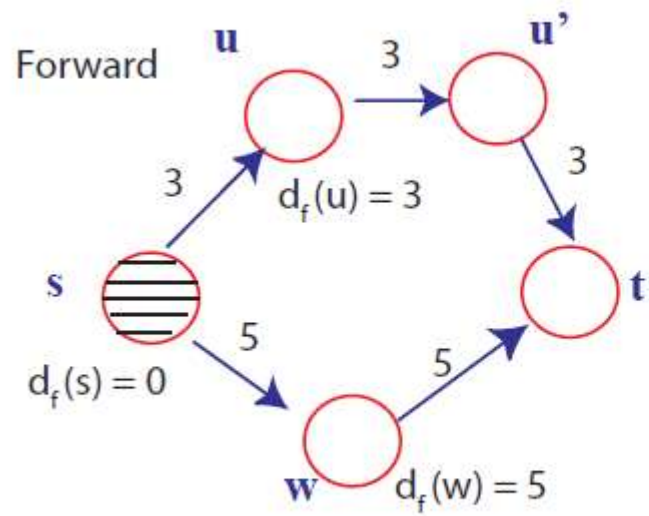
Minimum value for $d_f(x) + d_b(x)$ over all vertices that have been processed in at least one search

$$d_f(u) + d_b(u) = 3 + 6 = 9$$

$$d_f(u') + d_b(u') = 6 + 3 = 9$$

$$d_f(w) + d_b(w) = 5 + 5 = 10$$





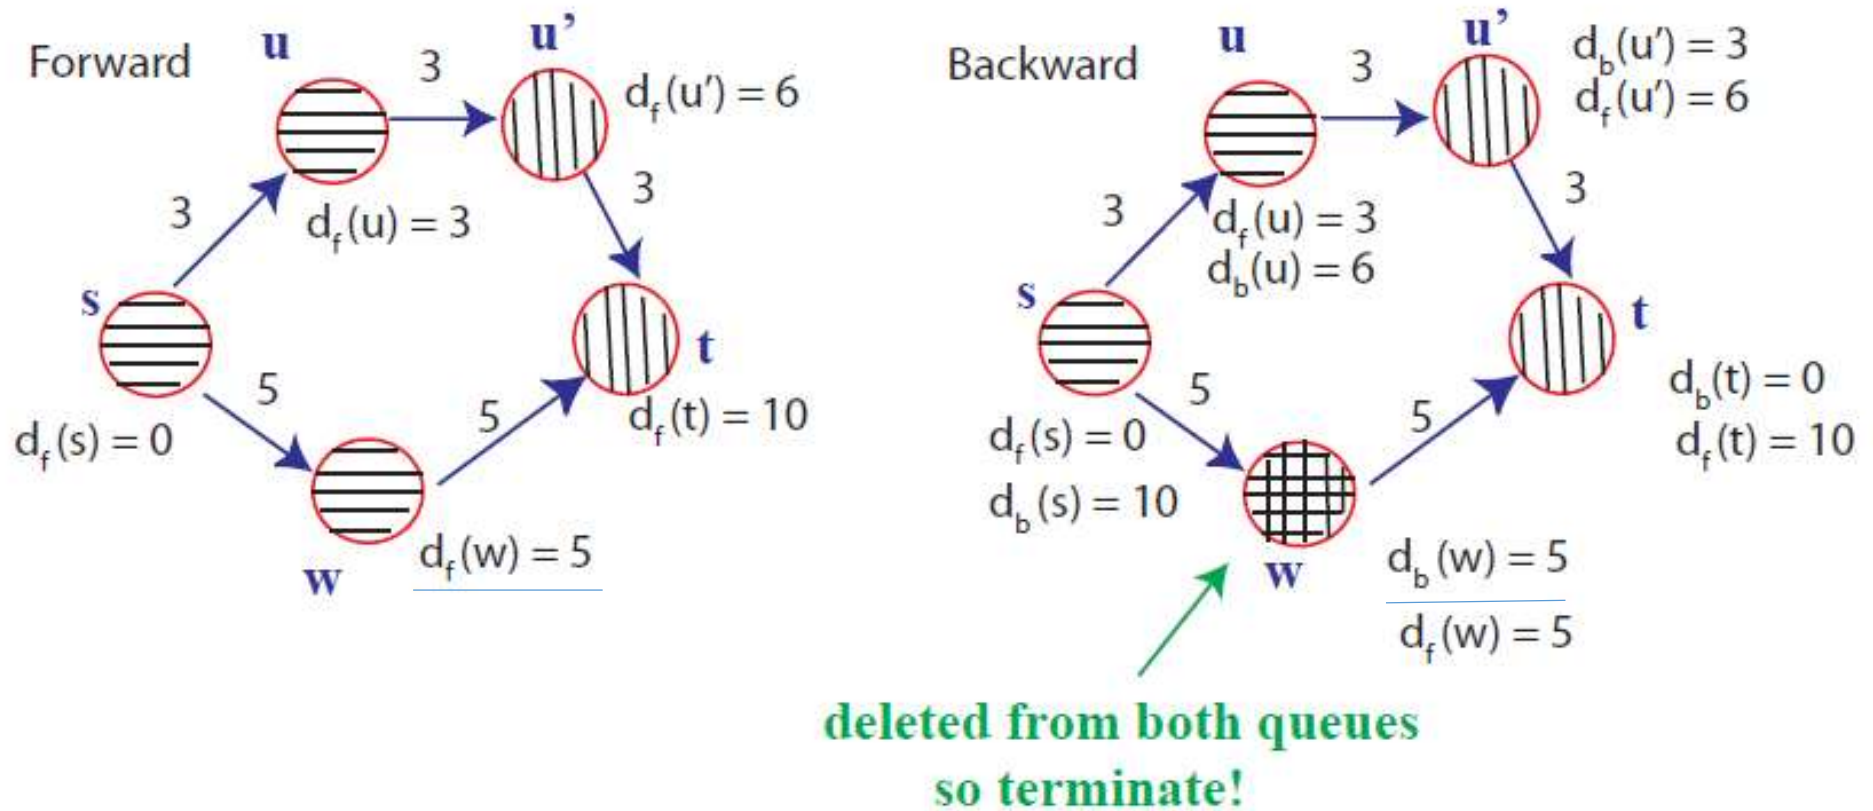


Figure 3: Forward and Backward Search and Termination.