

# Chapter 6: CPU Scheduling

---





# Background

---

- ❑ CPU scheduling is the basis of multi-programming operating systems.
- ❑ By switching the CPU among processes, the operating system can make the computer more productive
- ❑ In this chapter, we introduce basic CPU-scheduling concepts and present several CPU-scheduling algorithms





# Background

---

- ❑ Review the life cycle of a process in a single-processor
  - A process is executed until it must wait, typically for the completion of some I/O request
  - In a simple computer system, **CPU time is wasted!**
- ❑ Objective of multiprogramming
  - Having some process running at all times to maximize CPU utilization
  - **When one has to wait, takes the CPU away, gives to another**

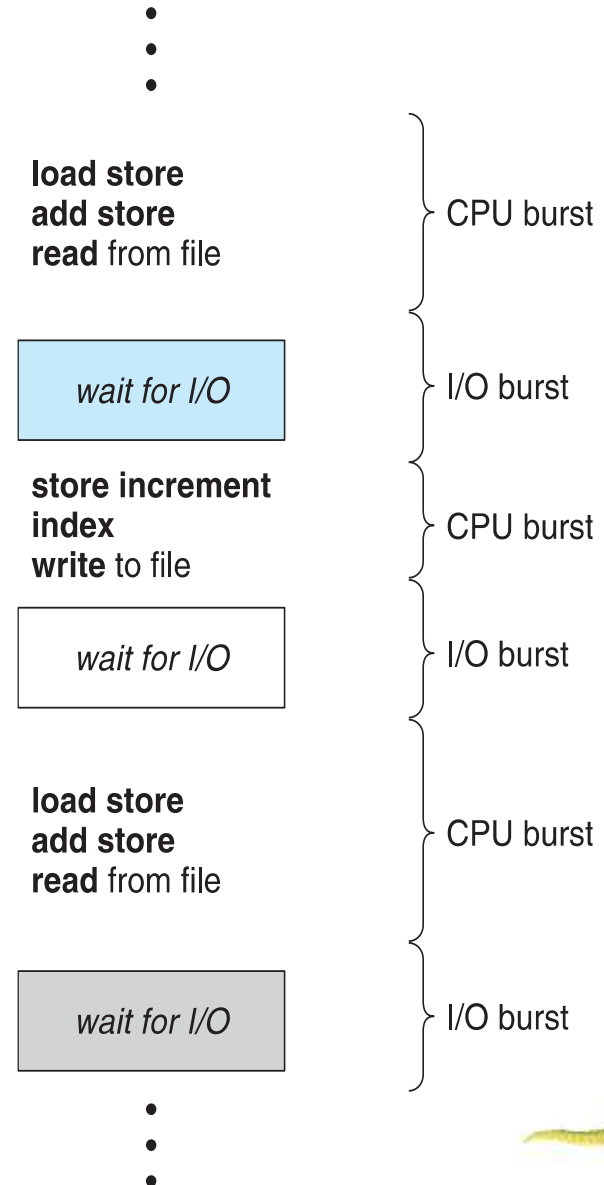
## **CPU scheduling**





# Basic Concepts

- ❑ CPU–I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❑ **CPU burst** followed by **I/O burst**
- ❑ Maximum CPU utilization obtained with multiprogramming
- ❑ CPU burst distribution is of main concern





# Review: Schedulers

---

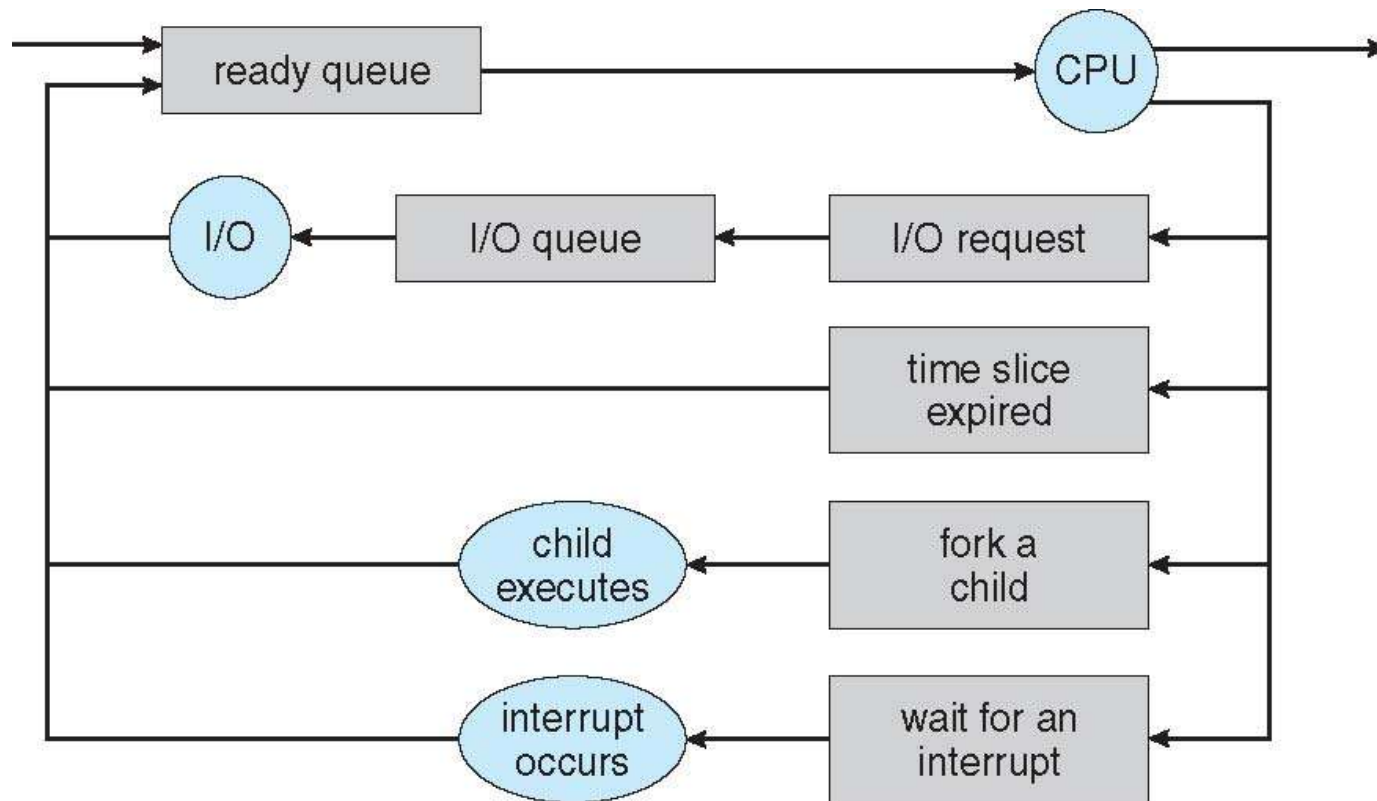
- ❑ **Short-term scheduler** (or **CPU scheduler**) – selects which process should be executed next and allocates CPU
  - Short-term scheduler is invoked **frequently** (milliseconds)
  - It must be fast, because of the short time between executions
  
- ❑ **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
  - Long-term scheduler is invoked **infrequently** (seconds, minutes), may be slow
  - The long-term scheduler controls the **degree of multiprogramming**





# CPU Scheduler

- ❑ **Short-term scheduler** (or **CPU scheduler**) – selects from among the processes in ready queue, and allocates CPU to one of them
  - Queue may be ordered in various ways
- ❑ **When can process scheduling take place?**





# Scheduling Criteria

- ❑ What criteria make sense?
- ❑ **CPU utilization** – keep the CPU as busy as possible
- ❑ **Throughput** – # of processes that complete their execution per time unit
- ❑ **Turnaround time** – amount of time to execute a particular process: from **time of submission** to **time of completion**
- ❑ **Waiting time** – amount of time a process has been waiting in the ready queue
  - CPU scheduling algorithm affects only the amount of time that a process spends waiting in the ready queue
- ❑ **Response time** – amount of time it takes from when a request was **submitted** until the **first response** is produced, **not output** (for time-sharing environment)

Difference between Turnaround time and Response time?





# Scheduling Algorithm Optimization Criteria

---

- ❑ Max CPU utilization
- ❑ Max throughput
- ❑ Min turnaround time
- ❑ Min waiting time
- ❑ Min response time







# First- Come, First-Served (FCFS) Scheduling

FCFS: the process that requests the CPU first is allocated the CPU first

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$






# FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

❑ The Gantt chart for the schedule is:



- ❑ Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- ❑ Average waiting time:  $(6 + 0 + 3)/3 = 3$
- ❑ Much better than previous case
- ❑ The average waiting time under the FCFS policy is often quite long
- ❑ The FCFS scheduling algorithm is non-preemptive

Thus, the average waiting time under the FCFS policy is generally **not minimal** and may vary substantially if the processes' CPU burst times vary greatly





# Shortest-Job-First (SJF) Scheduling

---

- ❑ Associate with each process the length of its **next CPU burst**
  - Use these lengths to schedule the process with the shortest time
  - When the CPU is available, it is assigned to the process that has the ***smallest next CPU burst***
  - See example in next slide
  
- ❑ Moving a short process before a long one decreases the waiting time of the short process more than it increases the waiting time of the long process. Consequently, the average waiting time decreases






# Example of SJF

<u>Process</u>	<u>Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

## ❑ SJF scheduling chart



❑ Average waiting time =  $(3 + 16 + 9 + 0) / 4 = 7$

❑ If in FCFS, Average waiting time =  $(0 + 6 + 14 + 21) / 4 = 10.25$





# Shortest-Job-First (SJF) Scheduling

---

- ❑ SJF is **optimal** – gives **minimum average waiting** time for a given set of processes
  - The difficulty is knowing the length of the next CPU request
  - Could ask the user





# Shortest-remaining-time-first

---

- ❑ The SJF algorithm can be either preemptive or non-preemptive
  - Non-preemptive: Once the CPU has been allocated to a process, that process keeps the CPU until it releases the CPU, either by terminating or by requesting I/O
  
- ❑ **Preemptive SJF** scheduling is also called **shortest-remaining-time-first**
  - When a new process arrives at the ready queue while a previous process is still executing
  - The next CPU burst of the newly arrived process may be **shorter than what is left of the currently executing process**
  - A preemptive SJF algorithm will preempt the currently executing process






# Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

- Preemptive* SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$  msec

Non-preemptive SJF would result in an average waiting time of 7.75 msec





# Priority Scheduling

- ❑ A priority number (integer) is associated with each process
- ❑ The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
    - When a process arrives at the ready queue, its priority is compared with the priority of the currently running process.
    - A preemptive priority scheduling preempt the CPU if the priority of the newly arrived process is higher than the priority of the currently running process.
  - Non-preemptive
- ❑ SJF is priority scheduling where priority is the inverse of next CPU burst time
- ❑ See example in next slide








# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

## ❑ Priority scheduling Gantt Chart



## ❑ Average waiting time = 8.2 msec





# Priority Scheduling

---

- ❑ Problem  $\equiv$  **Starvation** – low priority processes may never execute in preemptive case
- ❑ Solution  $\equiv$  **Aging** – gradually increasing the priority of processes that wait in the system for a long time





# Round Robin (RR)

---

- ❑ Similar to FCFS, but **preemption** is added to enable the system to switch between processes
- ❑ Each process gets **a small unit of CPU time** (**time quantum**  $q$ ), usually 10-100 milliseconds. **After this time has elapsed, the process is preempted and added to the end of the ready queue.**
- ❑ Treat the ready queue as a FIFO queue of processes. New processes are added to the tail of the ready queue. The CPU scheduler picks the first process from the ready queue, **sets a timer to interrupt after 1 time quantum**, and dispatches the process
- ❑ **Timer interrupts every quantum to schedule next process**





# Round Robin (RR)

---

- ❑ One of two things will then happen
- ❑ If the process have a CPU burst of less than 1 time quantum
  - Release the CPU voluntarily
  - The scheduler will then proceed to the next process in the ready queue
- ❑ If the CPU burst of the currently running process is longer than 1 time quantum
  - The timer will go off and will cause an interrupt to the operating system.
  - A context switch will be executed, and the process will be put at the tail of the ready queue.
  - The CPU scheduler will then select the next process in the ready queue.






# Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- ❑ Arrive at time 0, time quantum = 4
- ❑ The Gantt chart is:




- ❑ Average waiting time:  $[(10-4) + (4-0) + (7-0)] / 3 = 17/3 = 5.66$
- ❑ Typically, higher average turnaround than SJF, but better **response**





# Time Quantum and Context Switch Time

- ❑ The performance of the RR algorithm depends heavily on the size of the time quantum.
  - $q$  large  $\Rightarrow$  FCFS
  - $q$  small  $\Rightarrow$  results in a large number of context switch
  - $q$  should be large compared to context switch time, otherwise overhead is too high





# Multilevel Queue

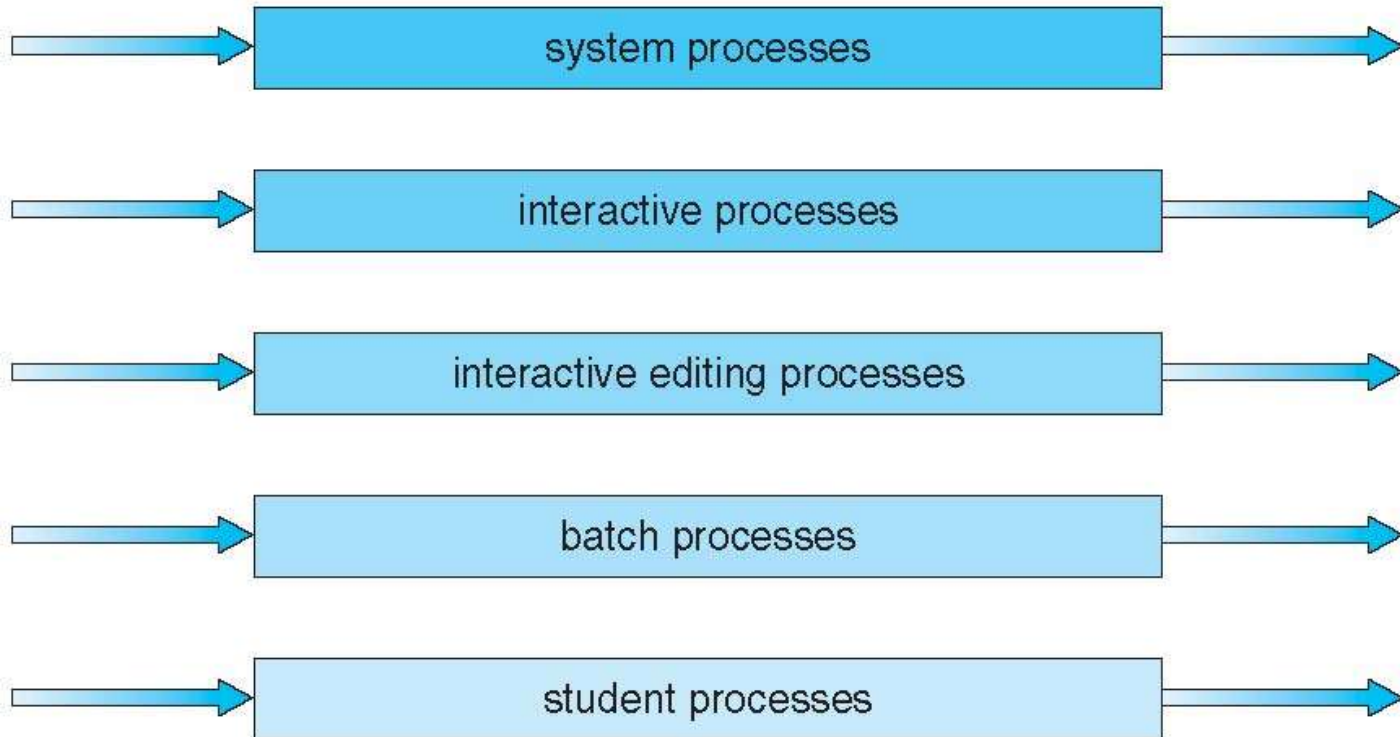
- ❑ A **multilevel queue** scheduling algorithm partitions the ready queue into several separate queues
- ❑ The processes are permanently assigned to one queue,
  - generally based on some property of the process, such as memory size, process priority, or process type.
- ❑ Each queue has its own scheduling algorithm:
  - foreground (interactive) – RR
  - background (batch) – FCFS
- ❑ Scheduling must be done between the queues:
  - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation. (**See Example**)
  - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR, and 20% to background in FCFS





# Multilevel Queue Scheduling

highest priority



lowest priority







# Multilevel Feedback Queue

---

- ❑ Allow a process to **move** between the various queues
- ❑ Idea: separate processes according to the characteristic of their **CPU bursts**
  - If a process uses too much CPU time, it will be moved to a lower-priority queue
  - Leaves I/O-bound and interactive processes in the higher-priority queues
  - In addition, a process that waits too long in a lower-priority queue may be moved to a higher-priority queue
    - ▶ This form **aging** prevents starvation






# Example of Multilevel Feedback Queue

## ❑ Three queues:

- $Q_0$  – RR with time quantum 8 milliseconds
- $Q_1$  – RR time quantum 16 milliseconds
- $Q_2$  – FCFS

## ❑ Scheduling

- A new job enters queue  $Q_0$  which is served FCFS
  - ▶ When it gains CPU, job receives 8 milliseconds
  - ▶ If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$
- If queue  $Q_0$  is empty, at  $Q_1$  job is again served FCFS and receives 16 additional milliseconds
  - ▶ If it still does not complete, it is preempted and moved to queue  $Q_2$





# Priority-based Scheduling with preemption

---


- ❑ For real-time scheduling, scheduler must support preemptive, priority-based scheduling
  - But only guarantees **soft** real-time functionality
- ❑ A process currently running on the CPU will be **preempted** if a **higher-priority** process available to run
- ❑ For **hard** real-time must also provide ability to meet **deadlines**
- ❑ **Deadline**: by which it must be serviced by the CPU





# Process Characteristics

- ❑ Processes have characteristics: **periodic**----- process requires CPU at constant intervals
  - Has fixed processing time  $t$ , deadline  $d$ , period  $p$
  - $0 \leq t \leq d \leq p$
  - **Rate** of periodic task is  $1/p$





# Rate Monotonic Scheduling

- ❑ The rate-monotonic scheduling algorithm schedules **periodic** tasks using a **static priority policy with preemption**.
- ❑ If a lower-priority process is running and a higher-priority process becomes available to run, it will preempt the lower-priority process.
- ❑ A priority is assigned based on the inverse of its period
  - Shorter periods = higher priority;
  - Longer periods = lower priority
  - Rationale: assign a higher priority to tasks that require the CPU more often





## Example: Using Rate Monotonic Scheduling

- Suppose  $P_1$  is assigned a higher priority than  $P_2$  because the period of  $P_1$  is shorter than that of  $P_2$ .




Figure 6.17 Rate-monotonic scheduling.

Deadline is satisfied!

Period:  $P_1 = 50$ ;  $P_2 = 100$ ;

Processing time:  $t_1 = 20$ ;  $t_2 = 35$ ;

Deadline: complete its CPU burst by the start of its next period.





## Example: Without Using Rate Monotonic Scheduling

- Suppose  $P_2$  is assigned a higher priority than  $P_1$




Figure 6.16 Scheduling of tasks when  $P_2$  has a higher priority than  $P_1$ .

P1 misses its first deadline at time 50

Period:  $P_1 = 50$ ;  $P_2 = 100$ ;

Processing time:  $t_1 = 20$ ;  $t_2 = 35$ ;

Deadline: complete its CPU burst by the start of its next period.





# Earliest Deadline First Scheduling (EDF)

---

- ❑ Priorities are **dynamically** assigned according to **deadlines**:
  - the earlier the deadline, the higher the priority;
  - the later the deadline, the lower the priority
  
- ❑ What the difference between rate-monotonic scheduling and EDF?
  - In EDF, priorities may have to be adjusted to reflect the deadline of the newly runnable process.
  - However, in rate-monotonic scheduling, priorities are fixed







## Earliest Deadline First Scheduling (EDF) (cont.)

### Example: Earliest Deadline First Scheduling (EDF)




Figure 6.19 Earliest-deadline-first scheduling.

- At time 50,  $P_2$  now has a higher priority than  $P_1$  because its next deadline (at time 80) is earlier than that of  $P_1$  (at time 100).
- $P_2$  is preempted at time 100 because  $P_1$  has an earlier deadline (time 150) than  $P_2$  (time 160).

Period:  $P_1 = 50$ ;  $P_2 = 80$ ;

Processing time:  $t_1 = 25$ ;  $t_2 = 35$ ;

Deadline: complete its CPU burst by the start of its next period.





## Problem 01 - Question

---

Consider three process, all arriving **at time zero**, with total execution time of **10, 20 and 30** units respectively.

Each process spends the first **20%** of execution time doing **I/O**, the next **70%** of time doing **computation**, and the last **10%** of time doing **I/O** again.

The operating system uses a **shortest-remaining-time-first** scheduling algorithm and schedules a new process either when the running process gets blocked on I/O or when the running process finishes its compute burst.

Assume that all I/O operations can be **overlapped** as much as possible.

For what percentage of does the CPU remain idle?






# Problem 01 - Solution

According to question, we have

	Total Burst Time	I/O Burst	CPU Burst	I/O Burst
Process P1	10	2	7	1
Process P2	20	4	14	2
Process P3	30	6	21	3

The scheduling algorithm used is Shortest-Remaining-Time-First:



Percentage of time CPU remains idle  
 $= (5 / 47) \times 100\%$   
 $= 10.638\%$






## Problem 02 - Question

Consider the set of 4 processes whose arrival time and burst time are given below

Process No.	Arrival Time	Burst Time		
		CPU Burst	I/O Burst	CPU Burst
P1	0	3	2	2
P2	0	2	4	1
P3	2	1	3	2
P4	5	2	2	1


If the CPU scheduling policy is **Shortest Remaining Time First**, calculate the **average waiting time** and **average turn around time**.






# Problem 02 - Solution

Process No.	Arrival Time	Burst Time		
		CPU Burst	I/O Burst	CPU Burst
P1	0	3	2	2
P2	0	2	4	1
P3	2	1	3	2
P4	5	2	2	1





# Problem 02 - Solution



Turn Around time = Exit time – Arrival time

Waiting time = Turn Around time – Burst time

$$\text{Average Turn Around time} = (11 + 7 + 7 + 11) / 4 = 36 / 4 = 9$$

$$\text{Average waiting time} = (6 + 4 + 4 + 8) / 4 = 22 / 4 = 5.5$$

Process Id	Exit time	Turn Around time	Waiting time
P1	11	$11 - 0 = 11$	$11 - (3+2) = 6$
P2	7	$7 - 0 = 7$	$7 - (2+1) = 4$
P3	9	$9 - 2 = 7$	$7 - (1+2) = 4$
P4	16	$16 - 5 = 11$	$11 - (2+1) = 8$





## Problem 03 - Question

Consider the set of 4 processes whose arrival time and burst time are given below

Process No.	Arrival Time	Priority	Burst Time		
			CPU Burst	I/O Burst	CPU Burst
P1	0	2	1	5	3
P2	2	3	3	3	1
P3	3	1	2	3	1


If the CPU scheduling policy is Priority Scheduling, calculate the average waiting time and average turn around time. (Lower number means higher priority)





# Problem 03 - Solution


Process No.	Arrival Time	Priority	Burst Time		
			CPU Burst	I/O Burst	CPU Burst
P1	0	2	1	5	3
P2	2	3	3	3	1
P3	3	1	2	3	1







## Problem 03 - Solution



Turn Around time = Exit time – Arrival time

Waiting time = Turn Around time – Burst time

$$\text{Average Turn Around time} = (10 + 13 + 6) / 3 = 29 / 3 = 9.67$$

$$\text{Average waiting time} = (6 + 9 + 3) / 3 = 18 / 3 = 6$$

Process Id	Exit time	Turn Around time	Waiting time
P1	10	$10 - 0 = 10$	$10 - (1+3) = 6$
P2	15	$15 - 2 = 13$	$13 - (3+1) = 9$
P3	9	$9 - 3 = 6$	$6 - (2+1) = 3$



# End of Chapter 6

---

